

Bug Report

Defects raised from the Playwright regression cycles of 03 and 05 October 2026. Twenty-three entries: three reproduced with screenshot, video and trace; sixteen held under a failing regression case that watches for the fix; the rest raised from behaviour observed repeatedly during the automation build.

1. Document control

Document	Bug Report — PMCH Healthcare ERP Test Automation
Version	2.0
Date raised	03 October 2026 (BUG-001 ... BUG-008), 05 October 2026 (BUG-009 ... BUG-023)
Raised by	QA Automation — EMSL Dev Team
Test cycle	Full regression run of 05 Oct 2026, 15:28 UTC — 99 tests in 33 m 33 s: 94 expected (78 passed and 16 expected-to-fail), 2 failed, 3 skipped. Supersedes the run of 03 Oct 2026, 08:04 UTC, which covered 24 of these cases.
Companion document	docs/TEST_CASES.pdf — the 110 cases these defects were found against. Each defect below names the case that holds it.
Evidence	docs/assets/failures/ — screenshots, videos, Playwright traces and page snapshots, with manifest.json describing the run. Regenerate with <code>npm run docs:evidence</code> .
Video playback	This PDF cannot play video. Open docs/src/bug-report.html in a browser for the inline players, or open the .webm files directly — paths are given with each defect and in the evidence index, section 7.

2. Environment

Application	HIS ERP — Healthcare Information System (Blazor Server / MudBlazor)
Tenant	Popular Medical College and Hospital Limited
Host	Hospital server over HTTPS, port 9006 (BASE_URL in .env). Live database.
Account	Administrator (credentials held in .env, not in this document)
Client	Chromium (Desktop Chrome profile), 1920×1080, headed, single worker
Harness	Playwright Test 1.63 · TypeScript · Windows 11

3. Legend

Severity	Critical — data loss or the module cannot be used · High — a core workflow produces a wrong result · Medium — workflow completes but behaves wrongly or needs a workaround · Low — cosmetic.
Priority	P1 fix before release · P2 fix this cycle · P3 backlog.

Evidence strength	REPRODUCED captured in a run with artifacts attached — screenshot, video, trace · HELD BY A CASE a regression case marked <code>test.fail()</code> asserts the correct behaviour and fails on every run; the run reports it as expected-to-fail, and will report it as <i>unexpectedly passing</i> the day the module is corrected, which is the signal to close the defect and drop the marker · OBSERVED seen repeatedly while building the suite and recorded in the test code, but not captured as an artifact.
--------------------------	---

4. Defect summary

ID	TITLE	MODULE	SEVERITY	PRI	EVIDENCE
BUG-001	Patient Registration list opens empty — Show From / Show To default to today	Registration	Medium	P2	REPRODUCED screenshot, 2 videos, trace
BUG-002	Advance payment taken in the admission wizard never reaches the Advance Dashboard	Admission / Accounts	High	P1	OBSERVED across runs
BUG-003	Brute-force throttle answers <i>every</i> request from the IP with HTTP 429	Authentication	Medium	P2	OBSERVED , measured
BUG-004	Patient Registration grid is slow to render its first row — 10 s typical, 34 s at worst	Registration	Medium	P2	REPRODUCED run timings, both cycles
BUG-005	Blazor circuit is dropped when more than one browser session is open	Platform	Medium	P2	OBSERVED
BUG-006	Every module route paints "Access Denied" for a third of a second before the real page	Platform	Medium	P2	OBSERVED
BUG-007	ERP host is served over HTTPS with a self-signed certificate	Infrastructure	Medium	P2	OBSERVED
BUG-008	Two misspelled buttons on the create-patient wizard: "Registration", "Previews"	Registration	Low	P3	OBSERVED
BUG-009	Picking another patient does not empty the cart — one patient's services are billed to the next	Consultancy Service	Critical	P1	HELD BY A CASE CS-12
BUG-010	Leaving the tab and returning clears the patient but keeps the cart — services against nobody	Consultancy Service	High	P1	HELD BY A CASE CS-13
BUG-011	The Rate box is editable, and the rate a nurse types is the rate the patient is billed	Consultancy Service	Critical	P1	HELD BY A CASE CS-14, CS-22
BUG-012	Committing a line resets Service Date, so a service performed earlier cannot be entered as such	Consultancy Service	Medium	P2	HELD BY A CASE CS-15
BUG-013	Service Change — the discount box — is read by nothing; the cart and the bill ignore it	Consultancy Service	Medium	P2	HELD BY A CASE CS-16
BUG-014	The blank first row of the Doctor Name lookup satisfies the required-field guard	Consultancy Service	High	P1	HELD BY A CASE CS-17
BUG-015	A lookup search that matches nothing shows a blank bordered row, not an empty state	Consultancy Service	Low	P3	HELD BY A CASE CS-18

ID	TITLE	MODULE	SEVERITY	PRI	EVIDENCE
BUG-016	The bill grid stops at 15 rows with no pager — a 16th charge cannot be seen or cancelled	Consultancy Service	High	P1	HELD BY A CASE CS-23
BUG-017	Diet indent quantity is not recorded — every line is filed as 1, zero included	Diet Indent	High	P1	HELD BY A CASE DI-06, DI-10
BUG-018	Diet Dashboard opens filtered to <i>tomorrow</i> , so the ward's own meals are not on it	Diet Dashboard	High	P1	REPRODUCED 2 screenshots, 4 videos, 2 traces
BUG-019	The date boxes are blank while a range is in force — the filter is invisible	Diet Dashboard	Medium	P2	REPRODUCED screenshot
BUG-020	A refused range and a search that matched nothing both answer with a silent empty dashboard	Diet Dashboard	Medium	P2	HELD BY A CASE DD-06, DD-09
BUG-021	Each grid stops at 120 rows with no pager and no total; older indents are invisible	Diet Dashboard	High	P1	HELD BY A CASE DD-07
BUG-022	Neither grid is labelled, so whether a patient is on a diet is readable only from the button	Diet Dashboard	Medium	P2	HELD BY A CASE DD-12
BUG-023	Stop takes a patient off their diet on one click, and Continue cancels the running one, both unannounced	Diet Dashboard	High	P1	HELD BY A CASE DD-16, DD-17

One item has been closed. AUT-001, the automation defect that took the 03 Oct run down with BUG-001, is fixed and verified — see section 6. BUG-001 itself stands: the test working around the default filter does not make the default filter right.

Why so many at once. BUG-009 ... BUG-023 are the Nurse Station tabs — Consultancy Service, Diet Indent and Diet Dashboard — which were automated between the two cycles and had not been covered before. Fifteen of the sixteen cases marked `test.fail()` were written against behaviour found while building those suites; BUG-018 is the one genuinely new failure the 05 Oct run turned up.

5a. Defect detail — Registration, Admission, Platform

Raised in the 03 Oct cycle. All eight were re-checked against the 05 Oct run and all eight still stand; what changed is that BUG-001 no longer takes the suite down with it, because the automation defect behind that is fixed (section 6).

BUG-001 — Patient Registration list opens empty because the date filter defaults to today

Module / page	Registration > Patient Registration (/hospital/newregistration)
Severity / Priority	Medium / P2
Status	OPEN — reproduced on both attempts of the 03 Oct run. Still present on 05 Oct; TC_ADM_002 passed that day only because it now widens the range itself (AUT-001, section 6).
Found by	TC_ADM_002 — tests/regression/admission.regression.spec.ts:63
Reproducibility	Deterministic. Any day before the first patient of the day is registered — in the 03 Oct run, both attempts, 08:04 UTC.

STEPS TO REPRODUCE

1. Sign in as Administrator.
2. Click the **REGISTRATION** tile on the module launcher.
3. Click **PATIENT REGISTRATION** in the left drawer.
4. Look at the grid without touching the date filter.


EXPECTED RESULT

The patient list opens on a range that shows work in progress — the recent registrations — so a user arriving at the page can find and open a patient record without first reaching for the date pickers.

ACTUAL RESULT

Show From and **Show To** are both set to today (10/3/2026), so the grid shows only patients registered today. Before the day's first registration that is nobody: the headers render, the body is empty, and the pager reads **0-0 of 0**. The page gives no indication that a filter is responsible — there is no empty-state message and no hint that widening the range would bring rows back.

EVIDENCE



Screenshot — assets/failures/TC_ADM_002/screenshot-1.png, cropped to the filter bar and the grid; the file itself is the full page. Both date filters read **10/3/2026**, the grid body is empty under its six headers, and the pager reads **0-0 of 0**.

Video, attempt 1 — assets/failures/TC_ADM_002/video-1.webm (479 KB). The page loads, the grid paints its headers, and nothing arrives for the 30 seconds the test waits. Not playable in the PDF — open the HTML or the file itself.

Video, attempt 2 (retry) — assets/failures/TC_ADM_002/video-2.webm (525 KB). Identical behaviour on the retry, which is what rules out a transient load failure.

- **Playwright trace** — assets/failures/TC_ADM_002/trace-2.zip. Open with `npx playwright show-trace docs/assets/failures/TC_ADM_002/trace-2.zip` for the DOM snapshot, network log and console at each step.
- **Page snapshot** — assets/failures/TC_ADM_002/error-context-1.md and error-context-2.md. The accessibility tree at the moment of failure, including the empty row groups and the **0-0 of 0** pager.

ERROR AS THE RUN REPORTED IT

```
TimeoutError: locator.innerText: Timeout 30000ms exceeded.  
Call log:  
  - waiting for locator('table.mud-table-root')  
    .locator('tbody tr').first().locator('td').nth(1)  
  
  at tests/regression/admission.regression.spec.ts:32  
    const uhid = (await registrationPage.rows.first()  
      .locator('td').nth(1).innerText()).trim();
```

ANALYSIS

The grid did not fail to load — it loaded correctly and returned nothing, because the filter it loaded with asks for one day. The run itself proves the filter is the cause rather than the data: TC_REG_004, which asserts the same grid has at least one row, *passed* in the same run. It runs after TC_FLOW_001, which registers a patient — so by the time TC_REG_004 looked, there was one registration dated today and the same default filter returned it. TC_ADM_002 ran before that, and saw an empty day.

SUGGESTED FIX

- Default **Show From** to a range that covers recent work — the last 7 or 30 days — rather than today, with **Show To** staying at today.
- Failing that, render an empty state in the grid that names the filter: *"No patients registered between 10/3/2026 and 10/3/2026. Widen the date range to see earlier registrations."*

BUG-002 — Advance payment taken in the admission wizard never reaches the Advance Dashboard

Module / page	Hospital > New Admission (/hospital/patientadmission) → Hospital > Advance Dashboard
Severity / Priority	High / P1 — money is collected at the counter and does not appear against the admission
Status	OPEN — needs confirmation from the module owners
Found by	TC_FLOW_001 — tests/regression/admission.regression.spec.ts:121
Reproducibility	Every run of the end-to-end flow so far.

STEPS TO REPRODUCE

1. Sign in as Administrator and register a patient.
2. Open **New Admission** and search the new UHID.
3. Fill the **Contact Person** tab.
4. On **Advance Payment**, enter all three: card VISA / EBL POS **1000**, bKash **2000**, cash **3000**.
5. Confirm the **Total:** panel reads **6000**.
6. Complete **Admission Detail** and press **Create Patient Admission**.
7. Open **Hospital > Advance Dashboard** and look for the admission.

EXPECTED RESULT

The 6000 taken at admission is recorded against the new admission and shows on the Advance Dashboard, so the billing desk can see what the patient has already paid.

ACTUAL RESULT

The admission is created correctly — it appears on **Patient Dashboard For Billing Dept** with the right bed, ward, mobile and blood group — but the Advance Dashboard carries **no row** for it. The wizard's own Total panel had summed the three amounts to 6000 immediately before the submit, so the figures were accepted by the form and lost somewhere between the submit and the advance ledger.

EVIDENCE

- No screenshot: the run that produced this observation predates the evidence collector, and TC_FLOW_001 passes, so Playwright captures nothing for a passing test.
- The observation is recorded in the test itself, at tests/regression/admission.regression.spec.ts:189, as the reason the advance total is deliberately not asserted after the submit.
- The pre-submit total is asserted (Total: 6000), so every passing run is standing evidence that the amounts reached the form.

OPEN QUESTION FOR THE MODULE OWNERS

It is not yet established whether this is a defect or a missing step: **Advance Collection** is its own page, and it may be that the wizard's Advance Payment tab only records an intent and the money has to be receipted separately. The answer decides whether this is a code fix or a documentation and training issue — and either way the wizard should not present a Total that then goes nowhere.

SUGGESTED FIX

- If the wizard is meant to post the advance: fix the posting and add the Advance Dashboard assertion to TC_FLOW_001 so the fix stays fixed.

- If it is not: remove or relabel the Advance Payment tab so it does not read as money taken, and state where the receipt is actually raised.

BUG-003 — Brute-force throttle answers every request from the IP with HTTP 429, not just login attempts

Module / page	Authentication — affects the whole application
Severity / Priority	Medium / P2
Status	OPEN
Found by	The negative login cases (TC_LOGIN_003-007), whose after-effects reached unrelated pages.
Reproducibility	Consistent, after roughly seven failed logins from one address.

STEPS TO REPRODUCE

1. From one IP address, attempt to log in with a wrong password about seven times.
2. From the *same* IP, on a different machine or browser, request any page — including one an already-signed-in user is working on.

EXPECTED RESULT

Throttling protects the sign-in endpoint and the account being attacked. Other users on the same network, and other routes, keep working.

ACTUAL RESULT

The host answers **every** request from that address with **HTTP 429** and `Retry-After: 60`. Measured, it clears in about ten seconds rather than the sixty advertised. During the cool-off a Blazor page load does not degrade gracefully: the circuit never starts, so the page renders as dead chrome rather than showing a "too many attempts, try again shortly" message.

WHY IT MATTERS BEYOND THE TEST SUITE

A hospital floor is behind one NAT address. One person mistyping their password several times locks out *everyone* on that floor for the cool-off — reception, billing and the wards alike — with no message explaining why the ERP has stopped responding.

EVIDENCE

- Worked around in `utils/helpers.ts` — `gotoTolerant()` and `reloadTolerant()` retry past the 429, and their comments record the measured behaviour.
- `reloadTolerant()` also documents a second symptom: a plain reload during the cool-off does not surface the 429 at all, it throws `net:ERR_HTTP_RESPONSE_CODE_FAILURE`.

SUGGESTED FIX

- Scope the throttle to the sign-in endpoint and the account being attempted, not to every route from the IP.
- Make `Retry-After` truthful — advertise the ten seconds actually enforced.
- Serve a readable 429 page instead of letting the Blazor circuit fail silently.

BUG-004 — Patient Registration grid can take over 30 seconds to render its first row

Module / page	Registration > Patient Registration (/hospital/newregistration)
Severity / Priority	Medium / P2
Status	OPEN — INTERMITTENT — 33.8 s on 03 Oct, 10.3 s on 05 Oct. Slow either way, but not reproducibly over 30 s; see <i>Actual result</i> .
Found by	TC_REG_004 — tests/regression/registration.regression.spec.ts:71

STEPS TO REPRODUCE

1. Sign in as Administrator.
2. Open **Registration > Patient Registration**.
3. Time how long it takes for the first row to appear, on a range that has data.

EXPECTED RESULT

A single page of ten rows renders in a couple of seconds, like every other page in the suite.

ACTUAL RESULT

TC_REG_004 took **33.8 seconds** in the 03 Oct run, against a suite median of about 2 seconds. It does nothing but open the page, read six column headers and wait for one row — the time is the grid. The same slowness is why the project raises Playwright's default timeouts across the board: 90 s per test, 30 s per action, 45 s per navigation.

Re-measured on 05 Oct: 10.3 seconds. Three times faster, and still four to five times every other case on the same page in the same session. So the page is reliably slow and *sometimes* very slow, which points at load on the host or on the query rather than at a fixed cost. The defect stands, with the headline figure read as a worst case rather than as typical.

EVIDENCE

- Run timings in docs/assets/failures/manifest.json, and in the HTML report under playwright-report/. Comparison, 03 Oct run, same session: TC_REG_003 (same page, no row wait) 2.4 s; TC_REG_002 4.4 s; TC_REG_008 3.8 s.
- 05 Oct run, same session: TC_REG_004 10.3 s against TC_REG_003 2.7 s, TC_REG_002 4.3 s, TC_REG_008 4.3 s. The three comparison cases barely moved between the two runs; TC_REG_004 moved by 23 seconds.

SUGGESTED FIX

- Profile the query behind the grid — a 10-row page should not be fetching the whole patient table before paging it.
- Show a loading indicator while the fetch is in flight, so the page does not read as empty while it works.

BUG-005 — Blazor circuit is dropped when more than one browser session is open

Module / page	Platform — the Blazor Server circuit, all modules
Severity / Priority	Medium / P2
Status	OPEN
Found by	Suite configuration — the reason <code>workers: 1</code> is forced in <code>playwright.config.ts</code> .

STEPS TO REPRODUCE

1. Sign in as the same account in two or more browsers at once against this host.
2. Work in both.

EXPECTED RESULT

Each session keeps its own circuit. A second browser does not disturb the first.

ACTUAL RESULT

Sessions are dropped. Running the suite with more than one worker produced lost sessions and failures unrelated to what the tests asserted, which is why the whole suite is pinned to a single worker — and why the Firefox and WebKit projects are commented out rather than run alongside Chromium.

IMPACT

Beyond the suite: a user who opens the ERP on a second machine, or a shared account used at two counters, can knock the other session out. It also puts a ceiling on test throughput — the full suite is serial, so it runs in minutes where it could run in under one.

EVIDENCE

- `playwright.config.ts` — `workers: 1`, with the reason recorded in the comment above it.
- The Firefox and WebKit project blocks in the same file, commented out for the same reason.

SUGGESTED FIX

Review circuit and session handling on the server — in particular whether sessions are keyed by account rather than by circuit, and what the configured circuit retention limits are.

BUG-006 — Every module route paints "Access Denied" before the real page arrives

Module / page	Platform — every module route
Severity / Priority	Medium / P2
Status	OPEN
Found by	Page-object navigation — <code>BasePage.settleAuthorization()</code> in <code>pages/base.page.ts</code> .

STEPS TO REPRODUCE

1. Sign in as Administrator.
2. Navigate to any module route — for example `/hospital/newregistration`.
3. Watch the first half-second of the load.

EXPECTED RESULT

A user with permission never sees a denial. The route renders a loading state, then the page.

ACTUAL RESULT

The app prerenders its **NotAuthorized** template before the Blazor circuit connects, so **"Access Denied"** is on screen within about 50 ms of every module route and is swapped for the real page roughly a third of a second later. The denial is never real — but it is what the user sees first, on every navigation.

IMPACT

Visible to staff as a flash of a permissions error on every page change, which trains people to distrust the screen. It also makes a genuine permission failure indistinguishable from the false one, for a user and for a test: the suite had to be taught to wait for the real page rather than believe the denial, and a run of about one in twenty does come back genuinely denied and needs a second navigation.

EVIDENCE

- `pages/base.page.ts` — `settleAuthorization()`, whose comment records the measured 50 ms / 300 ms timings and the one-in-twenty genuine denial.

SUGGESTED FIX

Render an authorizing/loading state during prerender instead of the `NotAuthorized` template, and only show the denial once authorization has actually been evaluated.

BUG-007 — ERP host is served over HTTPS with a self-signed certificate

Module / page	Infrastructure — the whole host
Severity / Priority	Medium / P2
Status	OPEN
Found by	Suite configuration — <code>ignoreHTTPSErrors: true</code> is required in <code>playwright.config.ts</code> .

STEPS TO REPRODUCE

1. Open the ERP URL in a browser with no certificate exception configured.

EXPECTED RESULT

A certificate the browser trusts. No warning, no exception for the user to click through.

ACTUAL RESULT

The browser presents a certificate warning that the user has to bypass. The test suite has to set `ignoreHTTPSErrors` for the same reason.

IMPACT

Patient data is involved, and the warning is on a page where staff enter credentials. Training people to click through certificate warnings removes the one signal that would tell them about a real interception, and the suite's blanket `ignoreHTTPSErrors` would not catch one either.

SUGGESTED FIX

Issue a certificate from a CA the hospital's machines trust — an internal CA is enough for an internal host. Then drop `ignoreHTTPSErrors` from the config so the suite starts failing if the certificate goes wrong again.

BUG-008 — Two misspelled buttons on the create-patient wizard

Module / page	Registration > Add New — create-patient wizard (/hospital/patients/new)
Severity / Priority	Low / P3
Status	OPEN
Found by	Page object — pages/registration.page.ts:219-222, where both locators match the app's spelling.

STEPS TO REPRODUCE

- 1. Open **Registration > Patient Registration > Add New**.
- 2. Read the wizard footer buttons.

EXPECTED / ACTUAL

EXPECTED LABEL	ACTUAL LABEL	WHERE
Confirm Registration	Confirm Registration	Wizard footer — the primary action, and the most-read label in the module
Previous	Previews	Wizard footer — back navigation between the three tabs

"Previews" is the worse of the two: it is not a misspelling of anything, it is a different word, and it sits next to *Next* where a user is reading quickly.

NOTE FOR WHOEVER FIXES IT

The automation locators currently match both misspellings on purpose, so that they match the live buttons. Correcting the labels will break `pages/registration.page.ts` and `TC_FLOW_001` — please flag the change so the locators are updated in the same release.

5b. Defect detail — Nurse Station > Consultancy Service

All eight were found while automating the tab for the 05 Oct cycle. Each is held by a regression case marked `test.fail()`: the case asserts the behaviour the tab is supposed to have, fails on every run, and will report as *unexpectedly passing* the day the module is corrected. The route is `/hospital/nurse-station > Consultancy Service` throughout, and the spec is `tests/regression/consultancy-service.regression.spec.ts`.

BUG-009 — Picking another patient does not empty the cart, so one patient's services are billed to the next

Module / page	Nurse Station > Consultancy Service
Severity / Priority	Critical / P1
Status	OPEN — the worst of the eight on this tab
Found by	CS-12 — <code>consultancy-service.regression.spec.ts:515</code>
Reproducibility	Deterministic — every run, every pair of beds.

STEPS TO REPRODUCE

1. Open Nurse Station > Consultancy Service.
2. Click the patient in bed **A** in the ward list.
3. Enter one service — pick a service, a doctor, a schedule, press **Enter** on Qty. The line lands in the cart and Total Amount shows its value.
4. Without pressing Save, click the patient in bed **B**.

EXPECTED RESULT

Picking a patient starts that patient's bill. The cart is emptied and Total Amount returns to zero, because nothing in it belongs to the patient now on screen.

ACTUAL RESULT

The patient panel moves to B — name, bed and Inv.No all change — and the cart does not. A's services are still listed, and A's money is still in Total Amount, under B's name. **Press Save from here and B is billed for A's consultants.** Nothing on screen marks the lines as belonging to someone else, and the only cue that anything is wrong is that the nurse remembers entering them for the previous bed.

IMPACT

One mis-click in a ward list produces a wrong bill against a real patient, silently. It is also the easiest mistake on the tab to make: the ward list and the entry row are on the same screen, and moving down the beds is the normal way to work.

SUGGESTED FIX

Clear the cart, Total Amount and the Service Change box whenever the selected admission changes — the same reset the save already performs. If carrying lines over is ever wanted, it must be an explicit choice with the patient named in the prompt, not the default.

BUG-010 — Leaving the tab and coming back clears the patient but keeps the cart

Module / page	Nurse Station > Consultancy Service
Severity / Priority	High / P1
Status	OPEN
Found by	CS-13 — consultancy-service.regression.spec.ts:568
Related	BUG-009 — the same cart, not cleared on a different trigger

STEPS TO REPRODUCE

1. Pick a patient and enter one service into the cart.
2. Click away to **Medicine Indent**.
3. Click back to **Consultancy Service** and wait for the pane to settle.

EXPECTED RESULT

A cart with no patient against it should not exist. Either both survive the round trip, or neither does.

ACTUAL RESULT

The round trip *does* clear the patient panel — which is the half the tab gets right — and leaves the cart holding the lines, with a figure under them in Total Amount, against nobody at all. The next nurse to pick a patient inherits them, which is BUG-009 by another route.

NOTE FOR WHOEVER FIXES IT

Returning to the tab repaints the grid empty for about a second before refilling it, so a quick look can suggest the cart was cleared when it was not. The regression case waits for the pane to settle before reading it, for exactly that reason.

SUGGESTED FIX

Clear the cart with the patient panel. One reset, run on both triggers.

BUG-011 — The Rate box is editable, and the rate a nurse types is the rate the patient is billed

Module / page	Nurse Station > Consultancy Service
Severity / Priority	Critical / P1
Status	OPEN — confirmed through the save
Found by	CS-14 (at the form) — consultancy-service.regression.spec.ts:620 CS-22 (on the bill) — consultancy-service.regression.spec.ts:1108

STEPS TO REPRODUCE

1. Pick a patient, then pick a service the catalogue prices at, say, 1000. The Rate box fills with 1000.
2. Type over it — set it to **1**.
3. Add a doctor and a schedule and press Enter on Qty.
4. Press **Save** and read the row that lands on the bill.

EXPECTED RESULT

The rate belongs to the service catalogue. The box is read-only at the nurse station, and whatever is on screen, the line is costed at the catalogue rate. A ward nurse adding a consultant's visit is not the person who prices it.

ACTUAL RESULT

The box takes the new number, the cart line carries it, and **the saved bill row carries it too** — 1 instead of 1000, at amount 1. Nothing anywhere records that the catalogue said otherwise. Typed the other way, the same path overcharges.

IMPACT

Pricing is editable by every ward account, without an approval step or an audit trail, on a live bill. Both directions matter: revenue lost on an undercharge, and a patient overcharged with nothing on the record to explain it.

NOTE ON HOW THIS WAS TESTED

Both cases type the rate *down* to 1, never up. This is the hospital's live database: a rate typed up would be a charge the ward never meant to make. CS-22 cancels its own row off the bill afterwards, with a reason, like every other writing case.

SUGGESTED FIX

Disable the Rate box and cost the line server-side from the catalogue. If a negotiated rate is a real requirement, put it behind a permission and record who changed it, what from and what to.

BUG-012 — Committing a line puts Service Date back, so a service performed earlier cannot be entered as such

Module / page	Nurse Station > Consultancy Service
Severity / Priority	Medium / P2
Status	OPEN
Found by	CS-15 — consultancy-service.regression.spec.ts:663

STEPS TO REPRODUCE

1. Pick a patient and set **Service Date** to a date a few days back.
2. Pick a service, a doctor and a schedule. The back-date is still in the box.
3. Press Enter on Qty and watch the box.

EXPECTED RESULT

The date chosen is the date the service is recorded against, and it survives the commit.

ACTUAL RESULT

The box takes the date and keeps it through the whole of the entry — right up to the Enter — then drops it, a render after the row lands, back to the day the tab opened on. The saved row carries the moment of the save rather than either. A consultancy performed on the ward three days ago cannot be entered as having been performed three days ago.

SUGGESTED FIX

Carry the entry row's Service Date into the committed line and into the saved row, and stop resetting the box on commit. If the reset is deliberate, it should at least happen before the line is read, not after.

BUG-013 — Service Change, the discount box, is read by nothing

Module / page	Nurse Station > Consultancy Service
Severity / Priority	Medium / P2
Status	OPEN
Found by	CS-16 — consultancy-service.regression.spec.ts:720

STEPS TO REPRODUCE

1. Pick a patient and fill a line for a service that costs more than the discount.
2. Type a discount into **Service Change**, under Total Amount.
3. Press Enter on Qty and read the cart line.

EXPECTED RESULT

The discount comes off the amount, and the cart line's Service Change column carries it.

ACTUAL RESULT

The box keeps what was typed into it, so nothing on screen suggests it was ignored — but the cart line carries no Service Change and the amount is the full rate. The saved row does not carry it either. The form has the box, the cart has a column for it and the bill has a *Discount* column, and nothing reads any of the three.

SUGGESTED FIX

Either wire the box through to the line and the saved row, or remove it and the two columns. A control that silently does nothing is worse than no control: a nurse who types a discount has every reason to believe it was applied.

BUG-014 — The blank first row of the Doctor Name lookup satisfies the required-field guard

Module / page	Nurse Station > Consultancy Service
Severity / Priority	High / P1
Status	OPEN — services with no doctor are already on live bills
Found by	CS-17 — consultancy-service.regression.spec.ts:769

STEPS TO REPRODUCE

1. Pick a patient and pick a service.
2. Open the **Doctor Name** lookup without searching. The first row is a blank record — no code, no name.
3. Click that row. The form shows no doctor.
4. Press Enter on Qty.

EXPECTED RESULT

Doctor Name is required, and a blank record is not a doctor. The commit is refused with *Please select fields.* — which is exactly what happens if the lookup is never opened at all (CS-06).

ACTUAL RESULT

Having opened the lookup and clicked its blank row is enough to satisfy the guard. The line is committed, and saves, with no doctor against it. CS-11 lists the services already on live bills in that state.

IMPACT

A consultancy charge with no consultant named cannot be audited, queried or attributed. The required-field check is in place and gives a false sense that it is holding.

SUGGESTED FIX

Do not render an empty record in the lookup. Validate on the selected doctor's ID rather than on whether a selection event occurred.

BUG-015 — A lookup search that matches nothing shows a blank bordered row, not an empty state

Module / page	Nurse Station > Consultancy Service — Service and Doctor Name lookups
Severity / Priority	Low / P3
Status	OPEN
Found by	CS-18 — consultancy-service.regression.spec.ts:813

STEPS TO REPRODUCE

- 1. Pick a patient and search the Service lookup for a term no service matches.

EXPECTED RESULT

The panel says there is no such service.

ACTUAL RESULT

An empty bordered row with no text in it — the same picture as a lookup that has not finished loading. Nothing selectable comes back, which is correct; what is missing is the sentence saying so.

SUGGESTED FIX

Render the grid's empty-message row with text. The component already has the row — it is the content that is blank.

BUG-016 — The bill grid stops at fifteen rows, so a sixteenth charge cannot be seen or cancelled

Module / page	Nurse Station > Consultancy Service — the lower grid
Severity / Priority	High / P1
Status	OPEN — six of the first eighteen beds are already at the cap
Found by	CS-23 — consultancy-service.regression.spec.ts:844

STEPS TO REPRODUCE

1. Open Consultancy Service and step through the ward until a patient's bill grid shows fifteen rows.
2. Look for a pager, a page-size box, a total, or a "show more".
3. Save a further service against that patient and look for it on the bill.

EXPECTED RESULT

The bill lists every consultancy charged against the admission, or says that it does not.

ACTUAL RESULT

The grid stops at **fifteen**. There is no pager, no page-size box and no total anywhere on the tab. A sixteenth consultancy is charged and then cannot be seen here at all — and because the bin that opens the cancel dialog lives on the row, **it cannot be cancelled from this tab either**. On a long-stay patient the ward is working from a list that quietly stopped being complete.

IMPACT

Charges exist that the ward cannot see or reverse. It also bit the automation before it was understood: a writing case saved against a patient already at the cap and could not find its own row, which is why every writing case now checks the patient has room first.

SUGGESTED FIX

Page the grid, or at minimum show the total count and a warning when the list is truncated. The cancel action should be reachable for any charge on the admission, not only for the fifteen most recent.

5c. Defect detail — Nurse Station > Diet Indent and Diet Dashboard

Route `/hospital/nurse-station` > **Diet Indent** and **Diet Dashboard**. Specs: `tests/regression/diet-indent.regression.spec.ts` and `tests/regression/diet-dashboard.regression.spec.ts`. **BUG-018** is the one defect in this document that the 05 Oct run turned up as a genuine failure, with artifacts; the rest are held by cases marked `test.fail()`.

BUG-017 — The diet indent quantity is not recorded: every line is filed as 1, zero included

Module / page	Nurse Station > Diet Indent > Add Indent
Severity / Priority	High / P1
Status	OPEN
Found by	DI-06 — <code>diet-indent.regression.spec.ts:135</code> DI-10 records the same thing as a run annotation
Reproducibility	Deterministic — every line, every run.

STEPS TO REPRODUCE

1. Open Diet Indent > Add Indent and select a ward patient.
2. Pick a food pattern, type a quantity of **5**, press Enter.
3. Read the quantity in the indent table.
4. Repeat with **9**, and with **0**.
5. Save, then open Verify Indent and read the quantity there.

EXPECTED RESULT

The line carries the quantity typed, and a quantity of **0** is refused with a message naming the quantity — the way Medicine Indent refuses it (TC_NS_009).

ACTUAL RESULT

Every line lands in the indent table reading **1**, whatever was typed, and reaches Verify Indent as 1 as well. A quantity of 0 is taken like any other and the line is added — reading 1. The number the nurse typed is accepted, displayed in the box, and then discarded.

IMPACT

The canteen is told to serve one of everything. A ward ordering five of a pattern gets one, and nothing on either screen says the figure was changed. The zero case is the same defect seen from the other side: an order the nurse meant to cancel out becomes an order for one meal.

NOTE FOR WHOEVER FIXES IT

This is why every other case in the diet suite reads quantities back rather than assuming them, and why DI-09 checks the Verify Indent quantity against the indent table rather than against what was typed. When the fix lands, DI-06 will report as unexpectedly passing and those cases can tighten.

SUGGESTED FIX

Bind the quantity box to the line that is committed, and validate it above zero at the Enter, with the same message style the form uses for its other refusals.

BUG-018 — Diet Dashboard opens filtered to *tomorrow*, so the ward's own meals are not on it

Module / page	Nurse Station > Diet Dashboard
Severity / Priority	High / P1
Status	OPEN — new in the 05 Oct run, reproduced on both attempts of both cases
Found by	DD-02 — diet-dashboard.regression.spec.ts:153 DD-13 — diet-dashboard.regression.spec.ts:609
Reproducibility	Deterministic within the run: four attempts across the two cases, 15:45 – 15:50 UTC on 05 Oct 2026 (21:45 – 21:50 local), all four identical.

STEPS TO REPRODUCE

- 1. Open Nurse Station > **Diet Dashboard**.
- 2. Read the startDate and endDate values the page posts (they are in hidden inputs — see BUG-019).
- 3. Compare them against the client's date.
- 4. For DD-13: widen the range, go to another tab, come back, and read them again.

EXPECTED RESULT

The dashboard opens filtered to **today** at both ends, and returning to the tab puts it back on today. Today is the day the ward is serving meals on.

ACTUAL RESULT

Both ends come back as 2026-10-06 against a client date of 2026-10-05 — one day ahead. Nothing has been ordered for tomorrow, so both grids are empty and the ward's own running diets are not on screen. The same day-ahead value comes back after a tab round trip.

ERROR AS THE RUN REPORTED IT


```
Error: expect(locator).toHaveValue(expected) failed

Locator:   locator('input[name="StartDate"]')
Expected:  "2026-10-05"
Received:  "2026-10-06"
Timeout:   15000ms

- locator resolved to <input type="hidden" name="StartDate" value="2026-10-06"
  placeholder="Start Date" class="form-control flatpickr-input" />
  - unexpected value "2026-10-06"

at tests/regression/diet-dashboard.regression.spec.ts:162:52
```

EVIDENCE



Screenshot — assets/failures/DD-02/screenshot-1.png, cropped to the filter bar and the two grid headers; the file itself is the full page. Both date boxes show nothing but *Start Date* and *End Date* while the hidden inputs carry 2026-10-06, and both grids are rendered to bare headers. BUG-018 and BUG-019 in one picture.

Video, DD-02 attempt 1 — assets/failures/DD-02/video-1.webm (493 KB). The tab opens, both grids paint their headers, and the date value never becomes today. Not playable in the PDF — open the HTML or the file itself.

Video, DD-13 attempt 2 (retry) — assets/failures/DD-13/video-2.webm (1.0 MB). The range is widened to a week, the tab is left for Diet Indent and returned to, and the filter comes back to the same day-ahead default.

- **Playwright traces** — assets/failures/DD-02/trace-2.zip (5.0 MB) and assets/failures/DD-13/trace-2.zip (12.8 MB). Open with `npx playwright show-trace <file>` for DOM snapshots, network and console per step.
- **Page snapshots** — DD-02/error-context-1.md and -2.md, and the same pair under DD-13/. Each carries the accessibility tree at the moment of failure, with the hidden input and its value in it.
- Four attempts across the two cases, all four identical — which is what rules out a transient clock skew during the run.

ANALYSIS

The client machine, the browser and the test host all agreed on 05 Oct; the value is the server's. Either the application server's clock or time zone is a day ahead of the ward, or the default range is computed as *today + 1*. The first is the more likely and the more serious, because it would affect everything else the server date-stamps — saved indents, bill rows, the Service Date on a consultancy. Worth checking before treating this as a dashboard bug.

IMPACT

A nurse opening the dashboard sees an empty ward and has no way to tell that from a ward with nothing ordered — the tab says nothing either way (BUG-020), and the range in force is not on screen (BUG-019). Three regression cases — DD-03, DD-04 and DD-08 — skipped in this run for want of rows, which is this defect downstream.

SUGGESTED FIX

1. Check the application server's clock and time zone against the hospital's first. If it is ahead, that is the fix, and it is wider than this tab.
2. If the clock is right, correct the default range to today at both ends.
3. Either way, fix BUG-019 alongside it — the day a filter is wrong is the day it most needs to be visible.

BUG-019 — The date boxes are blank while a range is in force, so the filter is invisible

Module / page	Nurse Station > Diet Dashboard — the Start Date and End Date boxes
Severity / Priority	Medium / P2
Status	OPEN
Found by	DD-02 — diet-dashboard.regression.spec.ts:153
Related	BUG-018 — this is what makes it invisible to the ward

STEPS TO REPRODUCE

1. Open Diet Dashboard and look at the two date boxes without touching them.

EXPECTED RESULT

The boxes show the range the grids below them are filtered to.

ACTUAL RESULT

Placeholder text and nothing else — *Start Date*, *End Date* — while a real range is in force. The value lives in a hidden input the flatpickr never mirrors on arrival. A nurse who widens the range, leaves the tab and comes back has no way of telling what they are now looking at.

EVIDENCE

docs/assets/failures/DD-02/screenshot-1.png — the empty boxes above two empty grids, which is the whole of the problem in one picture.

SUGGESTED FIX

Initialise the visible picker from the same value the hidden input is posted with. Until then, no automated case can assert on what a nurse can actually read there — the suite reads the hidden inputs instead, and says so.

BUG-020 — A refused range and a search that matched nothing both answer with a silent empty dashboard

Module / page	Nurse Station > Diet Dashboard
Severity / Priority	Medium / P2
Status	OPEN
Found by	DD-06 (reversed range) — diet-dashboard.regression.spec.ts:295 DD-09 (no-match search) — diet-dashboard.regression.spec.ts:427

STEPS TO REPRODUCE

1. Set **End Date** to a week ago and **Start Date** to today — the range a nurse produces by filling the second box first — and press Show.
2. Separately, search for a patient who is not on the ward.

EXPECTED RESULT

A range the app cannot honour is refused, in the field or on screen. A search that matched nothing says it matched nothing.

ACTUAL RESULT

Both empty the two grids to bare headers and say nothing at all. That is also exactly what a genuinely quiet day looks like, and what BUG-018's wrong default looks like. Four different situations, one picture.

IMPACT

A nurse cannot tell a filter that matched nothing from one the app would not run, or from a ward that has ordered nothing. This is what made BUG-018 hard to see from the screen alone.

SUGGESTED FIX

Validate the range at the field and say so when it is backwards. Give both grids an empty-state row — *no diet indents in this range* — and a distinct one for a search that matched nothing.

BUG-021 — Each grid stops at 120 rows, with no pager and no total

Module / page	Nurse Station > Diet Dashboard — both grids
Severity / Priority	High / P1
Status	OPEN
Found by	DD-07 — diet-dashboard.regression.spec.ts:328
Related	BUG-016 — the same shape of defect on the Consultancy bill grid, at fifteen

STEPS TO REPRODUCE

1. Filter the dashboard to the last fortnight and count the rows in each grid.
2. Filter to the last four weeks and count again; compare the oldest indent in each.

EXPECTED RESULT

Doubling the window brings back something older. A ward busy enough to fill a grid over a fortnight was busy over the fortnight before it too.

ACTUAL RESULT

Both windows come back with exactly **120 rows a grid** and the same oldest indent in each. The list is cut at 120 and the rest is simply not there. There is no pager, no page-size box and no total anywhere on the tab, so nothing says so. Five days of this ward is already enough to fill both grids, which means the cut-off is reached in ordinary use rather than in some extreme case.

IMPACT

Any question about more than a few days of the ward's diets is answered with a silently truncated list. Sorting happens before the cut, so ascending and descending return different sets of rows — which is also why the regression case compares sort *direction* rather than comparing one list against the reverse of another.

SUGGESTED FIX

Page both grids, and show the total row count beside each.

BUG-022 — Neither grid is labelled, so whether a patient is on a diet is readable only from the button

Module / page	Nurse Station > Diet Dashboard — both grids
Severity / Priority	Medium / P2
Status	OPEN
Found by	DD-12 — diet-dashboard.regression.spec.ts:588

STEPS TO REPRODUCE

1. Open Diet Dashboard with rows in both grids.
2. Try to tell, from the page alone, which list a given row is in.

EXPECTED RESULT

Each grid is named — a caption, a heading over the card, or at minimum an accessible label. The upper grid is the diet each patient is currently on; the lower is everything else in the range.

ACTUAL RESULT

The two are rendered identically: same classes, the same nine headers, no caption, no heading above either card, no `aria-label`. They sit one above the other. Status does not separate them either — Pending, PR Submitted, Served and Reviewed rows appear in both. **The only difference is the button in the last cell:** Stop in the upper grid, Continue in the lower.

IMPACT

Whether the patient in front of you is on that diet *now* has to be inferred from a button label. Combined with BUG-023 — where that button acts on the first click — the cost of misreading which grid you are in is a patient taken off their meal.

SUGGESTED FIX

Put a heading over each grid: *Running diets* and *Stopped diets*, or whatever the ward calls them. A caption element would serve screen readers at the same time.

BUG-023 — Stop takes a patient off their diet on one click, and Continue cancels the running one, both unannounced

Module / page	Nurse Station > Diet Dashboard — the Stop and Continue buttons
Severity / Priority	High / P1
Status	OPEN
Found by	DD-16 (Stop) — diet-dashboard.regression.spec.ts:769 DD-17 (Continue) — diet-dashboard.regression.spec.ts:804
Related	BUG-022 — which grid a row is in is already hard to read

STEPS TO REPRODUCE

1. Find a patient in the upper grid and press **Stop**. Watch for a confirmation or a message.
2. For the second half: with two indents on one patient, press **Continue** on the older one and watch what happens to the newer.

EXPECTED RESULT

Stopping a patient's diet is a clinical action and is confirmed, or at least acknowledged afterwards. Continue, which displaces whatever the patient is on now, says that it is about to do so.

ACTUAL RESULT

Stop acts on the first click. Nothing is asked and nothing is shown: by the time anything could be read, the row has already moved to the lower grid and the status is already rewritten. **Continue** is the quiet half of the same gap — it does not only start the diet it was pressed on, it stops the one the patient is on, and that second effect is neither asked about nor reported. In both cases the row moving between two unlabelled grids is the entire feedback.

IMPACT

A patient can be taken off their meal by one mis-click, with no confirmation, no undo and no record on screen that it happened. The ward's own Diet Indent tab guards a far smaller thing — it refuses a repeated food pattern outright — which makes the asymmetry hard to defend.

NOTE ON HOW THIS WAS TESTED

Both cases file their own diet indent first and act only on that one, and both put the ward back before reporting — the assertion fails on purpose, and a failing test runs no further, so the restore has to come first.

SUGGESTED FIX

1. Confirm Stop, naming the patient and the diet.
2. Confirm Continue when it would displace a running diet, naming what is about to be stopped.
3. Acknowledge both afterwards, the way the indent forms acknowledge a save.

6. Automation defect (not an application defect)

AUT-001 — TC_ADM_002 assumes the registration grid is never empty

Where	tests/regression/admission.regression.spec.ts:29-34 (newestUhid())
Severity / Priority	Medium / P2 — it is why the 03 Oct run came back red
Status	FIXED — closed 05 Oct 2026. newestUhid() is now listedPatient(), which widens the window before reading the grid and reads the name and the UHID off one row in a single call. Verified by TC_ADM_002 passing in 13.2 s on the 05 Oct run, where it had previously timed out twice.
Related	BUG-001 (the application behaviour the test runs into) — still open

WHAT IS WRONG

newestUhid() opens the patient list and reads the UHID out of the first row without widening the date filter or checking that a row exists. On a day with no registrations yet the grid is legitimately empty (BUG-001), so the read waits the full 30 seconds and the test fails with a timeout that says nothing about what the case was meant to prove. The same unguarded assumption sits in TC_SMK_ADM_003.

WHY IT MATTERS

The failure reads as "the admission wizard is broken" when the wizard was never reached. It also makes the suite's result depend on the time of day it runs, which is the kind of flakiness that gets a red build ignored.

FIX, AS APPLIED

1. **DONE** listedPatient() calls registrationPage.showLastDays() before reading the grid — the same thing a user would do.
2. **DONE** The UHID is asserted to be ten digits or more, with the patient named in the failure message, so an empty grid now fails as an empty grid rather than as a thirty-second timeout.
3. **DONE** Name and UHID are read off the same row in one call, so a re-render cannot hand back two different patients.
4. **OPEN** **BUG-001 stays open.** The test working around the default filter does not make the default filter right.

7. Evidence index

Everything below is under docs/assets/failures/ and was written by `npm run docs:evidence`, which reads the Playwright HTML report from the last run. One folder per failed case, named by its case ID, plus `manifest.json` describing the run.

Two runs are represented. DD-02/ and DD-13/ are from the 05 Oct run and hold BUG-018 and BUG-019. TC_ADM_002/ is from the 03 Oct run and is kept because BUG-001 is still open — the collector was run with `--keep` so it was not cleared away. The manifest describes the 05 Oct run only, which is why it lists two failures and not three.

FILE	SIZE	WHAT IT SHOWS
DD-02/		– the dashboard opens filtered to today (BUG-018, BUG-019)

FILE	SIZE	WHAT IT SHOWS
DD-02/screenshot-1.png	158 KB	The tab as the nurse sees it: both grids rendered to bare headers, and both date boxes showing nothing but their placeholder text, while the hidden inputs carry 2026-10-06. BUG-018 and BUG-019 in one picture.
DD-02/screenshot-2.png	158 KB	The retry — the same picture. The behaviour is deterministic.
DD-02/video-1.webm · video-2.webm	493 KB · 461 KB	Both attempts, full recordings: navigation, the tab opening, the 15 s wait on the date value.
DD-02/trace-2.zip	5.0 MB	Playwright trace of the retry — DOM snapshots, network, console, per-step timings. <code>npx playwright show-trace <file></code>
DD-02/error-context-1.md · -2.md	10 KB each	Accessibility-tree snapshot of the page at failure, plus the failing source lines.
DD-13/ – the same default after a tab round trip (BUG-018)		
DD-13/screenshot-1.png · -2.png	158 KB each	The dashboard after leaving the tab and coming back: search cleared, range back to the day-ahead default.
DD-13/video-1.webm · video-2.webm	928 KB · 1.0 MB	Both attempts, including the trip out to Diet Indent and back — the nurse who goes to order a meal and returns.
DD-13/trace-2.zip	12.8 MB	Trace of the retry, with the full tab round trip step by step.
DD-13/error-context-1.md · -2.md	10 KB each	Page snapshot at failure and the failing source lines.
TC_ADM_002/ – kept from the 03 Oct run for BUG-001		
TC_ADM_002/screenshot-1.png · -2.png	52 KB each	The empty registration grid at the moment of failure, both date filters on 10/3/2026, pager 0-0 of 0. Byte-identical across the two attempts.
TC_ADM_002/video-1.webm · video-2.webm	479 KB · 525 KB	Both attempts: navigation, grid paint, the 30 s wait.
TC_ADM_002/trace-2.zip	1.7 MB	Trace of the retry.
TC_ADM_002/error-context-1.md · -2.md	12 KB each	Page snapshot at failure and the failing source lines.
manifest.json	6 KB	The 05 Oct run's statistics, per-attempt status and timings, and the error text, as machine-readable JSON.

The other fifteen defects have no artifacts, by design. Each is held by a case marked `test.fail()`, and Playwright records nothing for a test that failed as expected. The evidence for those is the case itself — its assertion, the comment above it, and the run annotations it writes, which the HTML report carries per test: what Save would have billed and to whom (CS-12), what the catalogue said against what the line charged (CS-14), what the quantity box was given against what the indent recorded (DI-10), and so on. Open `npm run report` and read the annotations under each case.

8. Recommended order of work

Ordered by what a defect costs the hospital if it is left alone, not by how hard it is to fix. The first four are all cases where a real patient record or a real bill ends up wrong, silently.

ORDER	ID	REASON
1	BUG-009, BUG-011	Both put a wrong charge on a real patient's bill with no warning and no audit trail — one patient's services billed to the next, and a rate any ward account can type over. Nothing else in this document is as cheap to trigger or as hard to detect after the fact. BUG-010 is the same cart, not cleared on a different trigger, so fix it in the same change.
2	BUG-018	Start with the server's clock and time zone, not with the dashboard. If the server is a day ahead, it is date-stamping saved indents and bill rows too, and that is a much wider problem than one empty grid.
3	BUG-002	Money is collected in the admission wizard and not recorded. Highest value of the older defects, and it needs a decision from the module owners before any code is written.
4	BUG-017, BUG-023	Both are patient care rather than money. A ward ordering five meals gets one, and a patient can be taken off their diet by a single mis-click with nothing asked and nothing said.
5	BUG-014, BUG-016, BUG-021	Records that cannot be audited or reached: charges with no consultant named, charges past the fifteenth that cannot be seen or cancelled, and diet history past 120 rows that is simply not there. BUG-016 and BUG-021 are the same defect in two grids and may share a fix.
6	BUG-001, BUG-019, BUG-020, BUG-022	The "what am I looking at" group — a list filtered to today with no sign of it, a range in force that is nowhere on screen, four different situations that all render as the same empty grid, and two grids that are not labelled. Individually small; together they are why BUG-018 was invisible from the screen alone. Worth one pass by one person.
7	BUG-003	One mistyped password can stall a whole floor behind the same NAT address.
8	BUG-006, BUG-004	Both are visible to staff on every working day — a false permissions error on every navigation, and a grid that takes half a minute.
9	BUG-012, BUG-013	Two controls on the Consultancy form that take input and then discard it: the back-date and the discount. Either wire them through or take them off the form.
10	BUG-007, BUG-005	Infrastructure and platform work; schedule with the next maintenance window.
11	BUG-008, BUG-015	Cosmetic. The two misspellings are the most-read labels in the registration module, and the blank lookup row is one sentence of text. Coordinate BUG-008 with the automation locators, which match both misspellings on purpose.

How to tell a fix landed. Sixteen of these defects are held by a regression case marked `test.fail()`. When the module is corrected, that case does not quietly start passing — the run reports it as **unexpectedly passing**, which fails the build until the marker is removed. So the suite will tell you, on the next run after the fix, and the defect can be closed against that run. Please flag such a change so the marker comes off in the same release.